

Lab 3 for uC/OS-II: Ceiling Priority Protocol

Prof. Li-Pin Chang

ESSLab@NCTU

Objective

- To implement Ceiling Priority Protocol for ucOS's mutex locks

uC/OS Mutex Locks

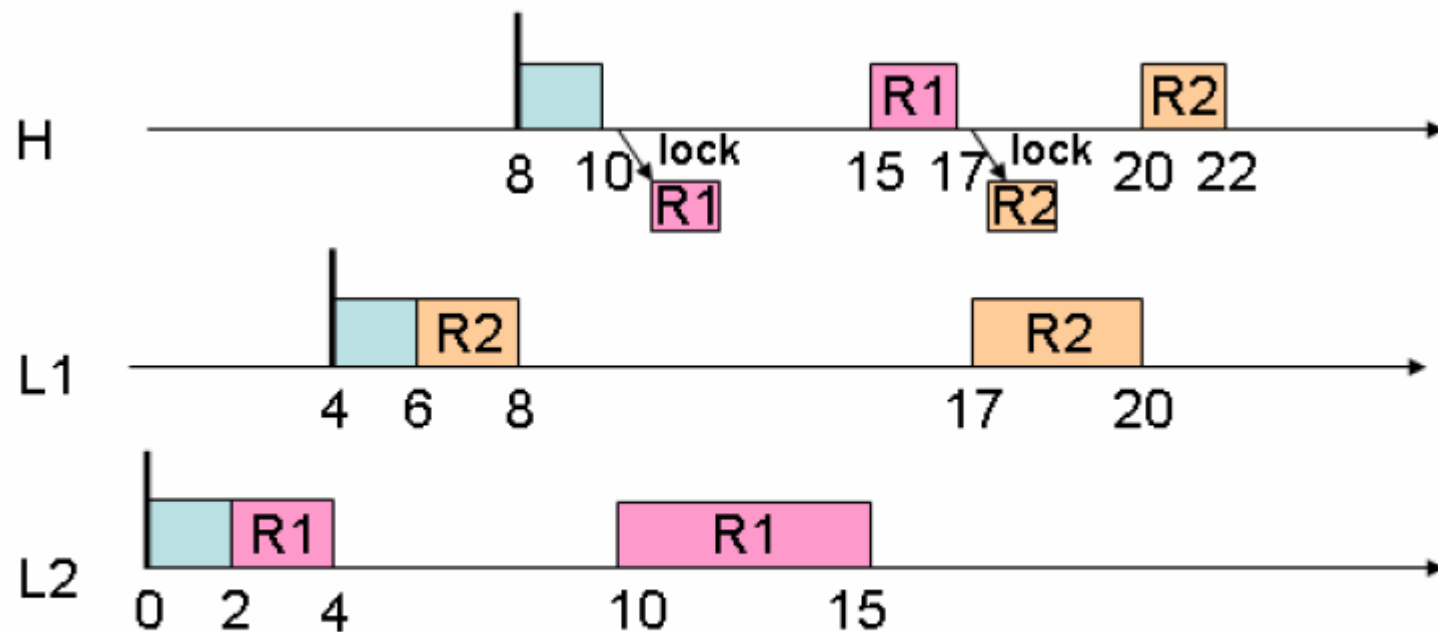
- A mutex lock is associated with a “priority”
 - Its priority is higher than the highest locker
 - E.g., if T_3 and T_4 share a lock, the priority of the lock should be set to 2
 - When T_4 **blocks** T_3 , then T_4 's priority becomes 2

Disadvantages of PIP

- The “PIP” avoids uncontrolled priority inversion, but it has two disadvantages
 - A high priority task can be blocked multiple times
 - Deadlocks are possible

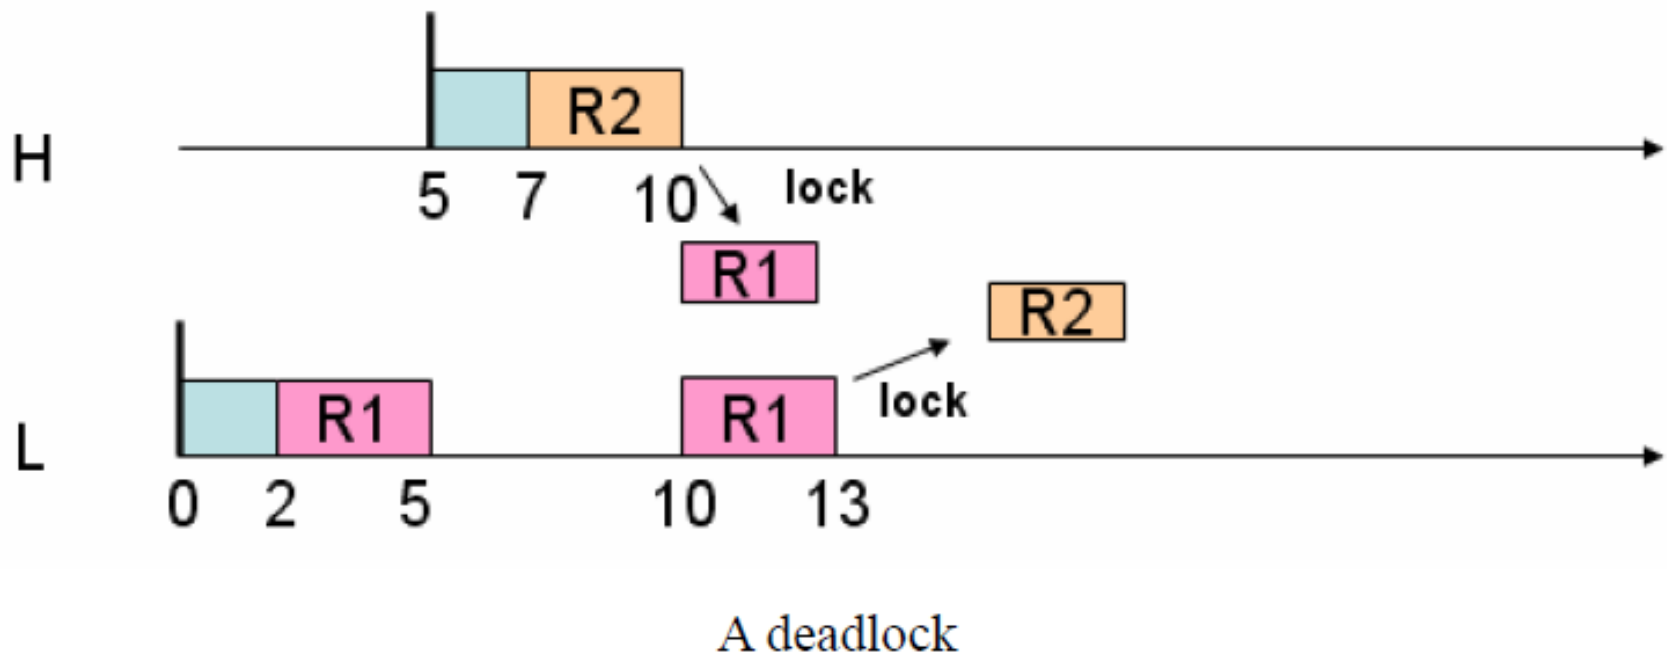
Scenario 1: Multiple blocking in ucOS2

PIP



Task H is in turn blocked by task L1 and task L2

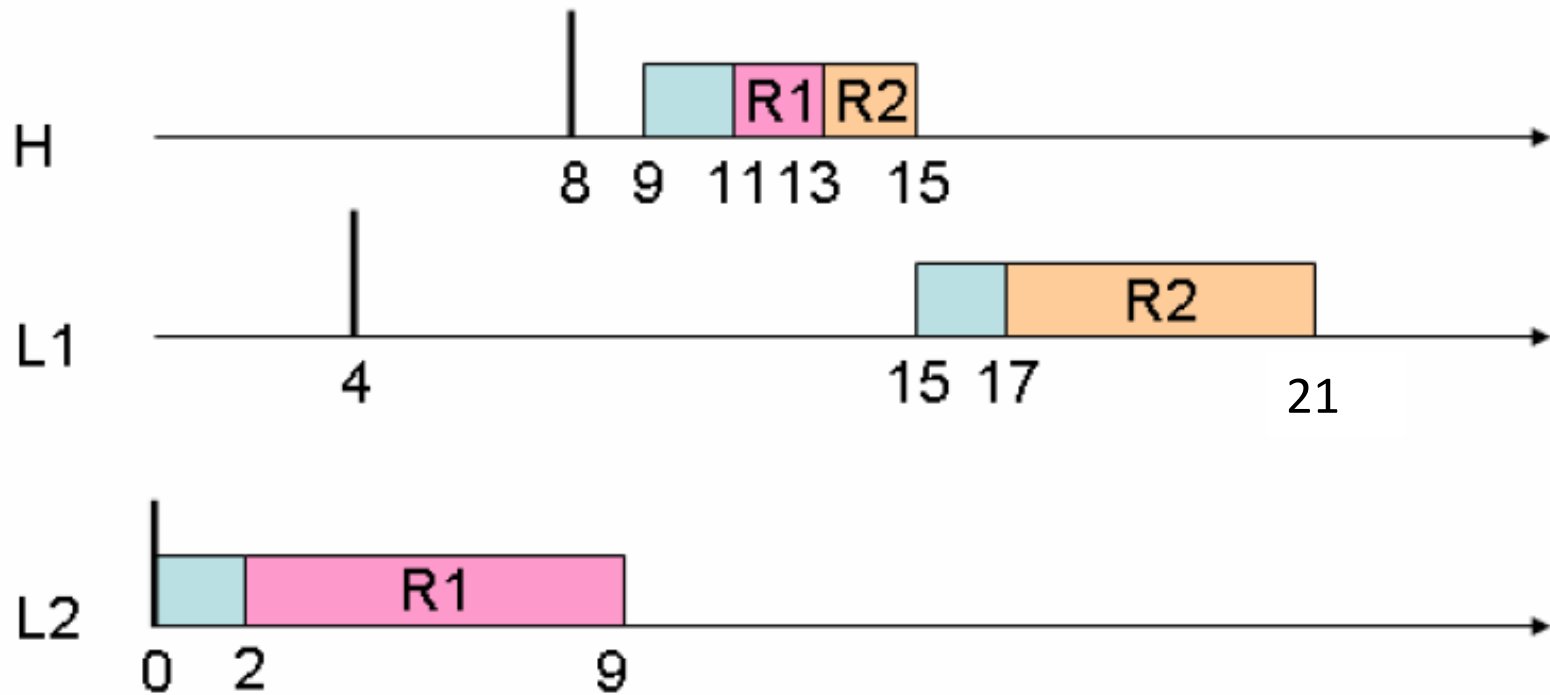
Scenario 2: Deadlock in uCOS-2 PIP



Ceiling Priority Protocol


- Highest-Locker Protocol
- When a task **acquires** a mutex lock, its priority becomes the highest among all lockers' priorities
- In uC/OS, the we use the mutex's priority as the highest-locker's priority
 - Immediately higher than all lockers' priorities

S1 CPP: Removing Multiple Blockings



The result of applying CPP

S2 CPP: Avoiding Deadlocks



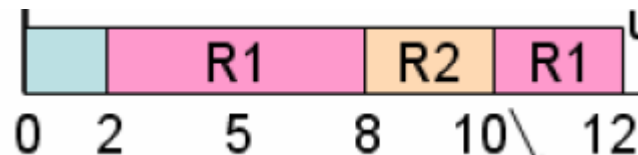
Deadlocks are avoided by using CPP

Implementation

- Reuse your code of Lab 1 (**do not re-use EDF**)
- Modify the following two functions
 - OSMutexPend()
 - If mutex is free, **boost** the locker's (caller) priority
 - OSMutexPost()
 - **Restore** the original priority of the locker
- **Do not** use OSTaskChangePrio()
 - It calls OS_Sched() and results in unexpected behaviors

Implementation

- All tasks should add proper OSTimeDly() at their beginning to emulate their arrival times
- Emulate durations of CPU execution and resource use with your code from Lab 0
 - 2 ticks → lock R1 → 6 ticks → lock R2 → 2 ticks → unlock R2 → 2 ticks → unlock R1



Output

- Similar to those in prior labs, but add lock/unlock events
- Output the results of using CPP for **Scenarios 1 and 2 (shown previously)**

Output Example of S1

Priority initialization:

R1: 1

R2: 2

Task1: 3

Task2: 4

Task3: 5

Task arrival time:

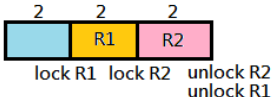
Task1: 8

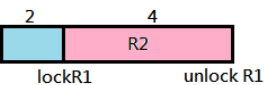
Task2: 4


Task3: 0


20	lock	R1	(Prio=5 changes to=1)
90	unlock	R1	(Prio=1 changes to=5)
90	complete		5 3
110	lock	R1	(Prio=3 changes to=1)
130	lock	R2	(Prio=1 changes to=1)
150	unlock	R2	(Prio=1 changes to=1)
150	unlock	R1	(Prio=1 changes to=3)
150	complete		3 4
170	lock	R2	(Prio=4 changes to=2)
210	unlock	R2	(Prio=2 changes to=4)
210	complete		4 19

Task execution time and resource used:

Task1: 

Task2: 

Task3: 



- Here, at time 90, it is actually a “preempt” because T3 calls OSMutexPost(R1), which internally calls OS_Sched() to surrender the CPU to T1 (HPT)
- Because our prior lab defines “OS_Sched()” as “complete” so it is okay. T3 ends after unlocking R1 anyway. The same to times 150 and 220.

20	lock	R1	(Prio=5 changes to=1)
90	unlock	R1	(Prio=1 changes to=5)
90	complete		5 3
<hr/>			
110	lock	R1	(Prio=3 changes to=1)
130	lock	R2	(Prio=1 changes to=1)
150	unlock	R2	(Prio=1 changes to=1)
150	unlock	R1	(Prio=1 changes to=3)
150	complete		3 4
170	lock	R2	(Prio=4 changes to=2)
220	unlock	R2	(Prio=2 changes to=4)
220	complete		4 19

